

# **verl: Flexible and Scalable Reinforcement Learning Library for LLM Reasoning and Tool-Calling**

Presenter: Haibin Lin, ByteDance Seed

# **Motivation: Why is Large-Scale RL Important?**

# Large-Scale RL for Reasoning and Agents

Learning to reason with large-scale RL greatly boosts the performance of LLMs

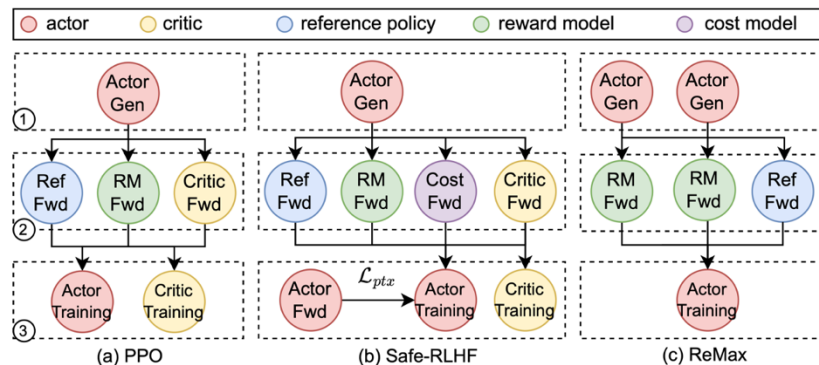
Model	Large-Scale RL?	AIME 2024	MATH 500	GPQA Diamond	Code Forces
GPT-4o ( <a href="#">OpenAI 2024</a> )	✗	44.6	60.3	50.6	>11.0%
o1 ( <a href="#">OpenAI 2024</a> )	✓	74.4	94.8	77.3	>89.0%

***Deep research ... was trained on real-world tasks requiring browser and Python tool use, using the same reinforcement learning methods behind OpenAI o1, our first reasoning model.***

– OpenAI Deep Research Blog, 2025

**Challenge:**  
**Why is Large-Scale RL**  
**Challenging?**

# RL as a Complex Dataflow

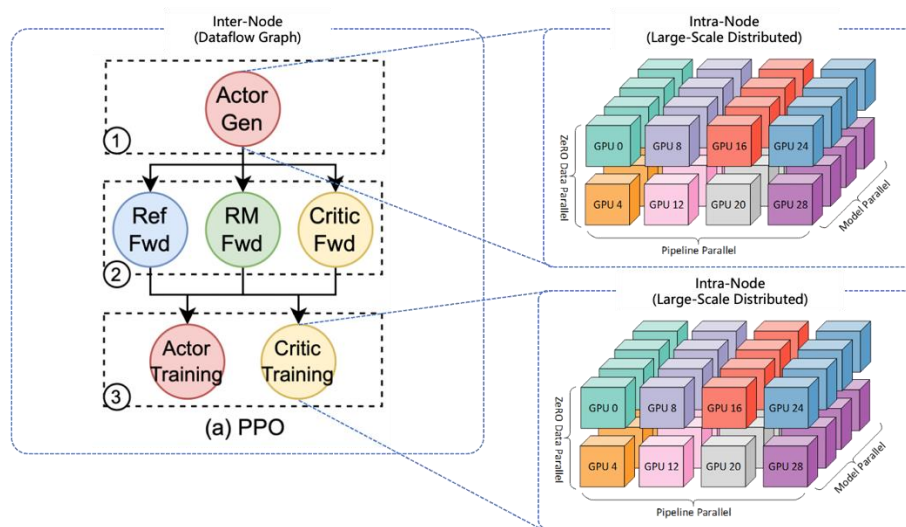


Reinforcement Learning (RL) can be modelled as **complex dataflow graph** ([Schaarschmidt et al. 2019](#); [Liang et al. 2021](#); [Sheng et al. 2025](#)), consisting of:

- **multiple models:** actor, critic, reference, reward model, etc.
- **multiple stages:** generating, preparing experiences, training
- **multiple workloads:** generation, inference, training

HybridFlow, Sheng et al., 2024

# RL with LLMs is Large-Scale Distributed Dataflow

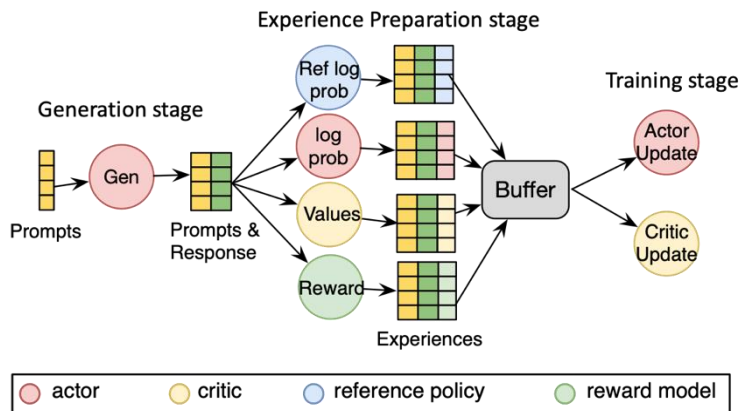


each **operator** in the RL dataflow = a large-scale **distributed** computing workload  
constraints: computation dependency

# Why verl for RL with LLMs?

Flexible and Efficient!

# Flexibility in Programming: “Single-Controller”



```
1 for prompts in dataloader:
2     # Stage 1: Generation
3     batch = actor.generate_sequences(prompts)
4     # Stage 2: Experience Preparation
5     batch = reward.compute_reward(batch)
6     batch = reference.compute_log_prob(batch)
7     batch = critic.compute_values(batch)
8     batch = compute_advantage(batch, "gae")
9     # Stage 3: Training
10    critic.update_critic(batch)
11    actor.update_actor(batch)
```

- Programming interface based on the “**single-controller**” paradigm
- RL algorithm core logic in **a few lines of code!**
- Diverse RL algorithms supported: PPO, GRPO, RLOO, GSPO, PRIME, DAPO, etc.



# Efficiency: “Multi-Controller”

verl is efficient for intra-operator with the “**multi-controller**” paradigm and features like:

## Training Backends:

- FSDP
- FSDP2
- Megatron

## Generation Backends:

- vLLM
- SGLang
- ...

## Parallelism Algorithms:

- Data Parallelism
- Tensor Parallelism
- Pipeline Parallelism
- Context / Sequence Parallelism
- ...


## Efficient Kernels:

- Flash Attention 2
- Torch Compile
- Liger Kernel
- ...

# Efficiency: “Hybrid Engine” for synchronous RL

verl is efficient for inter-operator with the “**hybrid engine**” paradigm, utilizing:

- **offloading & reloading** enables fully utilizing the GPU memory
- **resharding** enables switching for the optimal parallelism strategy
- FSDP trainer & vLLM worker live in the **same process**, reducing mem fragmentation



HybridFlow, Sheng et al., 2024

# Open-Source Community: Impactful and Inclusive

So far, verl has gained:

- 11.9k stars
- 2k forks
- 1.5k PRs
- 300+ contributors

Waiting for your participation!

Capabilities

- Multi-modal (image/video)
- Large MoE RL
- Multi-GPU LoRA support
- Sandbox/search tools
- Recipes: DAPO, retool




# Paradigm behind verl: HybridFlow

Sheng et al. 2024

# Background: Single-Controller vs. Multi-Controller


## Single-Controller (MPMD):

A centralized controller manages all the workers, running different programs



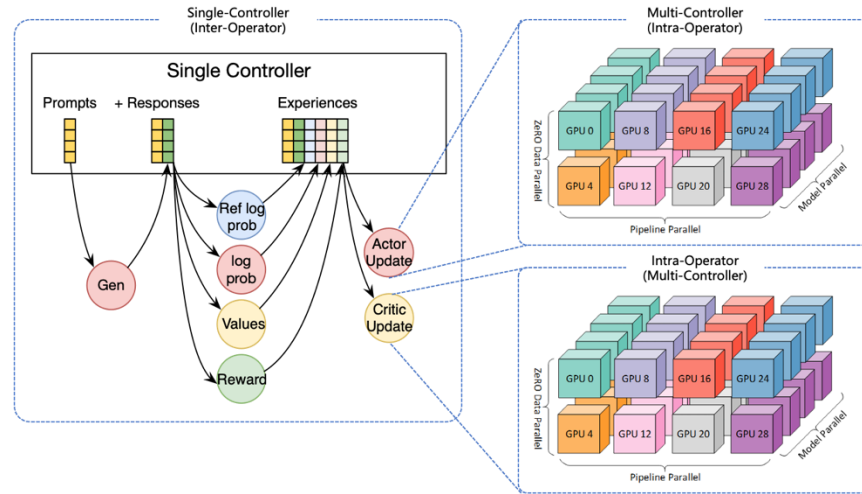
## Multi-Controller (SPMD):

Each worker has its own controller, running the same program with different data



Pathways, Barham et al., 2022

# New Paradigm: Hybrid-Controller!



- Hybrid-Controller = Single-Controller + N x Multi-Controller
- In the hybrid-controller, a single-controller manages multiple multi-controllers to process the dataflow


# Controller Data Dispatch/Collection in verl

```
1 for prompts in dataloader:
2     # Stage 1: Generation
3     batch = actor.generate_sequences(prompts)
4     # Stage 2: Experience Preparation
5     batch = reward.compute_reward(batch)
6     batch = reference.compute_log_prob(batch)
7     batch = critic.compute_values(batch)
8     batch = compute_advantage(batch, "gae")
9     # Stage 3: Training
10    critic.update_critic(batch)
11    actor.update_actor(batch)
```

```
1 class CriticWorker(3DParallelWorker):
2     @register(dispatch_mode=3D_PROTO)
3     def compute_values(self, batch: DataProto):
4         values = self.critic.forward(batch)
5         batch.update(values=values)
6     # ...
7 class ActorWorker(3DParallelWorker):
8     @register(dispatch_mode=3D_PROTO)
9     def update_actor(self, batch: DataProto):
10         loss = self.actor(batch)
11         loss.backward()
```

- Each call in the single-controller (e.g. `critic.compute_values`, `actor.update_actor`) is an RPC to a multi-controller worker group
- The register decorator utility manages the distributed data transfer, which also makes multi-controller programming easier

# Train/Generation Weight Transfer in verl



Training: TP=4 DP=2 PP=1

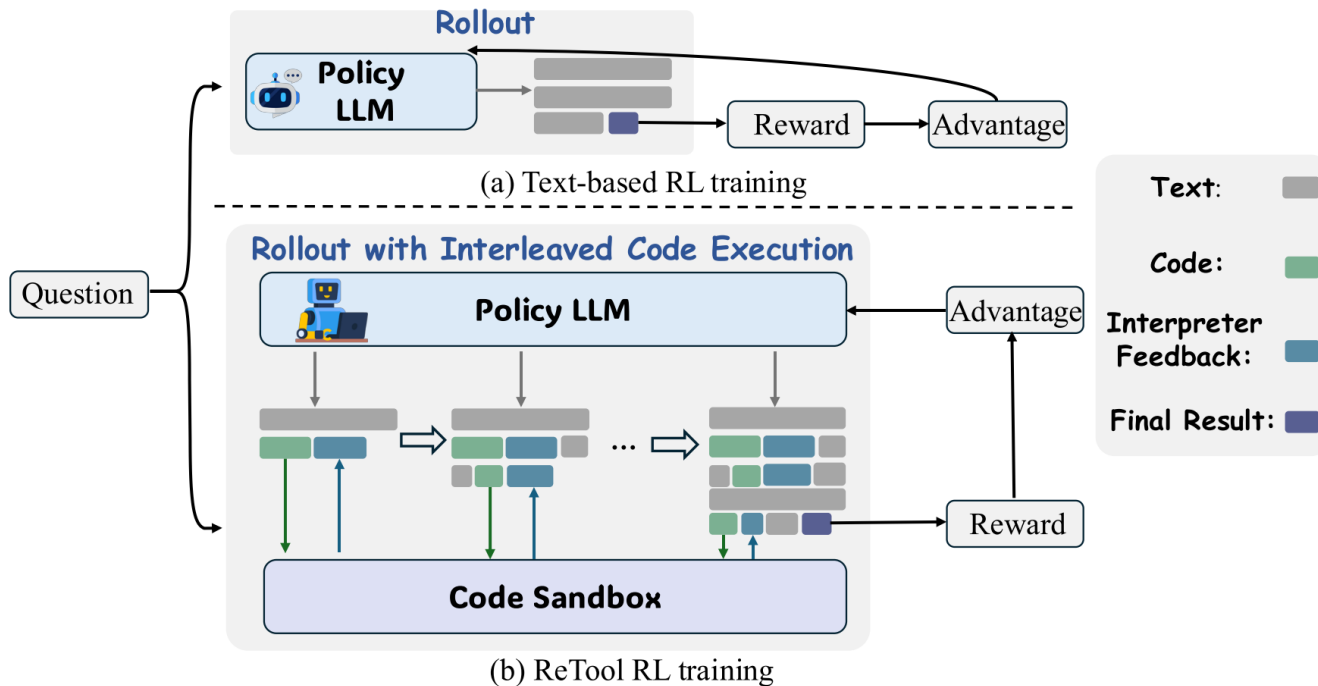
Generation: TP=2 DP=4 PP=1

- Optimal sharding for training/generation can be different
- Weight binding can be generalized as a transformation of DTensor between two device meshes under the same world via NCCL+RDMA
- Per-tensor/bucket transfer to avoid OOMs

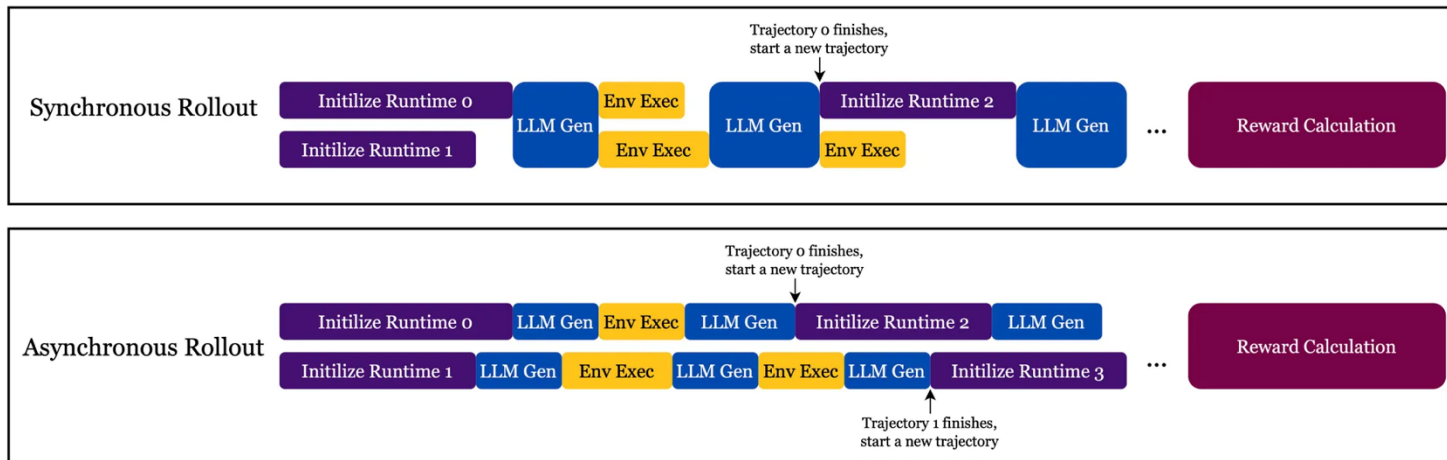
```
for shared_training_param in model.parameters():  
    train_full_param = shared_param.full_tensor()  
    infer_sharded_param = redistribute(train_full_param, infer_device_mesh)
```



# Approaching Agentic RL



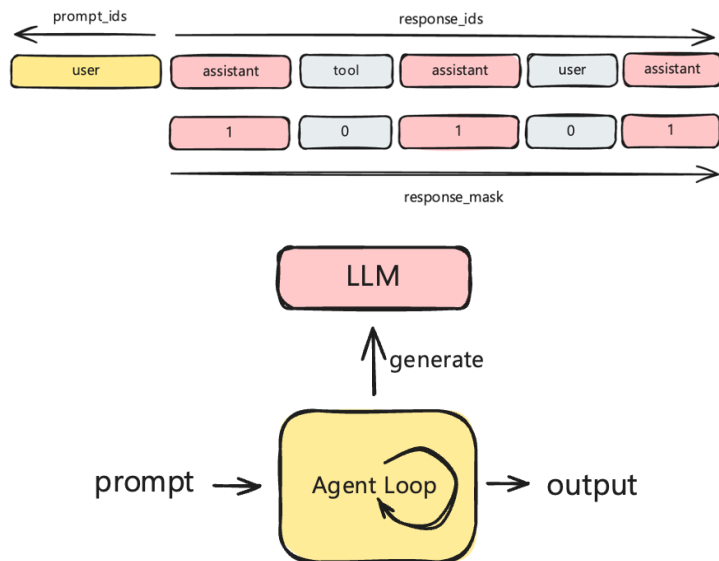
# Async Engine for Multi-Turn Rollout



- Synchronous Engine: returns all the outputs in the batch at the same time
- Asynchronous Server: returns each output as soon as it is ready

# Token-in-token-out Agent Loop Interface

Given one prompt, run a user defined loop with multi-turn/tool calling trajectories.  
Token ids are used for server generation input / output.



```
class DemoLoop(AgentLoopBase):  
  
    async def run(self, messages: list[dict[str, Any]], ...) -> AgentLoopOutput:  
        prompt_ids = await self.loop.run_in_executor(None,  
            lambda: tokenizer.apply_chat_template(messages, ..., tokenize=True)  
        )  
        num_turns = 0  
        while not is_done(prompt_ids, num_turns):  
            response_ids = await self.rollout_server.generate(  
                request_id=request_id, prompt_ids=prompt_ids, ...  
            )  
            prompt_ids += response_ids  
            tool_response_ids = await call_tool(response_ids)  
            prompt_ids += tool_response_ids  
            num_turns += 1  
            response_mask += ...  
  
        output = AgentLoopOutput(  
            prompt_ids=prompt_ids,  
            response_ids=response_ids,  
            response_mask=response_mask,  
            num_turns=num_turns,  
        )  
        return output
```

# Agent Loop & Reproducible Recipes

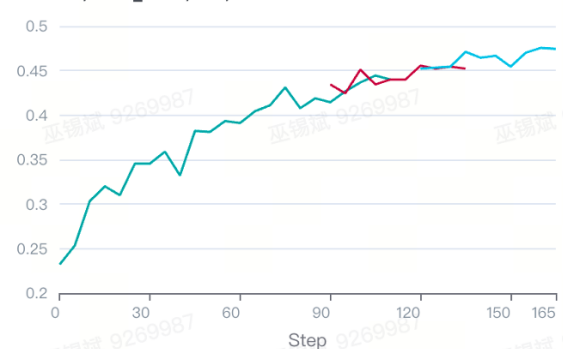
- client-server mode: decoupled inference engine and agent loop
- parallel async loop running: request/prompt level async execution
- KV cache reuse: multi-turn requests sent to the same inference server

## ReTool Recipe

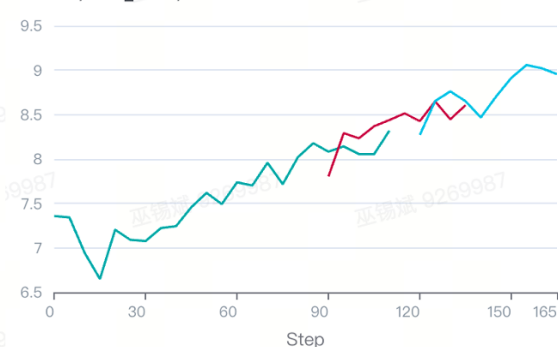
AIME 0.6 with 32b model

[verl/recipe/retool/](#)

val-core/aime\_2025/acc/mean@30



val-aux/num\_turns/mean



# **Recent Updates & Roadmap**

# Efficient RL with Huge MoE like DeepSeek-V3-671B

verl is working on supporting efficient RL training for **giant MoE like DeepSeek-V3-671B**, based on the following features:


- Runnable with 256 H100 GPUs
- Training: **MoE models** based on Megatron, ~0.12 MFU
- Inference: **Multi-node** tensor parallel inference
- Verified reward curve on orz57k & proprietary datasets from community
- Planned: accelerate rollout performance (e.g. fp8)

For more details, please check [issue tracker #1033](#).


# Disaggregated Async Trainer

- Hybrid controller allows flexible device placement (collocate/disaggregated)
- One-step async pipeline to avoid onload/offload overhead & bubbles
- [verl/recipe/one\\_step\\_off\\_policy](#)
- ~1.2x speedup compared to sync trainer
- Long-tail generation problems
- Towards fully-async pipeline


Interruptible generation [PR 2200](#)



(a) Disaggregated Architecture



(b) Colocated Architecture



(c) One-step Asynchronous Pipelining

# Q3 Roadmap

- Modular design: composable model engines with better abstraction
  - Algorithm agnostic engine abstraction: FSDP2, Megatron, and more
- Partial rollout & fully-async training pipeline ([AReal, Kimi, 2025](#))
- Rollout performance optimizations (fp8)
- Agentic RL recipes (e.g. SWE-bench)

Github roadmap [issue tracker #2388](#).

